

Coupon Recommendation System Using User Behavior Data

Project Description

- This project focuses on building a machine learning model to predict whether a user will accept a coupon based on several factors, such as weather, passenger information, time of day, and more. The dataset is drawn from an in-vehicle environment where recommendations are offered to users, and the goal is to make predictions that could improve coupon targeting for businesses.

Data Understanding

- The dataset used for this project is the In-Vehicle Coupon Recommendation dataset. It contains information about 12,684 users who were offered various types of coupons while driving. The dataset includes a variety of features related to user behavior and conditions in which the coupons were presented. Each row represents one instance of a coupon being presented, and the target variable (Y) indicates whether the user accepted the coupon or not.

Key Features:

- destination: Where the user is going (e.g., No Urgent Place, Work).
- passenger: Who the user is traveling with (e.g., Alone, Friend(s)).
- weather: The weather condition when the coupon was offered (e.g., Sunny, Rainy).
- temperature: The outside temperature in Fahrenheit.
- time: The time of day when the coupon was presented (e.g., 10AM, 2PM).
- coupon: The type of coupon offered (e.g., Coffee House, Restaurant(<20), etc.).
- expiration: The expiration time for the coupon (e.g., 2 hours, 1 day).
- has_children: Whether the user has children or not.
- Bar, CoffeeHouse, CarryAway, Restaurant: How often the user visits these establishments.
- direction_same: Whether the user is heading in the same direction as the coupon destination.
- Y: Target variable (1 for accepting the coupon, 0 for rejecting it).

Target Variable:

- Y: Indicates whether the coupon was accepted (1) or rejected (0).

Data Loading and Preview

- Before starting the analysis, it was important to set up a good set of tools. This included numpy and pandas for data handling, matplotlib and seaborn for creating visuals, and various features from sklearn for data processing, machine learning, and evaluating performance.

```
In [1]: from google.colab import drive
drive.mount('/content/drive')#Mount the drive to the colab env to access files

Mounted at /content/drive
```

```
In [2]: #importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score, roc_curve
```


```
In [3]: #Loading the data
df = pd.read_csv('/content/drive/MyDrive/in-vehicle-coupon-recommendation.csv')
```

```
In [4]: #view the first five rows of the data
df.head()
```

```
Out[4]:
```

	destination	passanger	weather	temperature	time	coupon	expiration	gender	age
0	No Urgent Place	Alone	Sunny	55	2PM	Restaurant(<20)	1d	Female	21
1	No Urgent Place	Friend(s)	Sunny	80	10AM	Coffee House	2h	Female	21
2	No Urgent Place	Friend(s)	Sunny	80	10AM	Carry out & Take away	2h	Female	21
3	No Urgent Place	Friend(s)	Sunny	80	2PM	Coffee House	2h	Female	21
4	No Urgent Place	Friend(s)	Sunny	80	2PM	Coffee House	1d	Female	21

5 rows × 26 columns




In [5]: *#view the last five columns of the data*
df.tail()

Out[5]:

	destination	passanger	weather	temperature	time	coupon	expiration	gender	ag
12679	Home	Partner	Rainy	55	6PM	Carry out & Take away	1d	Male	2
12680	Work	Alone	Rainy	55	7AM	Carry out & Take away	1d	Male	2
12681	Work	Alone	Snowy	30	7AM	Coffee House	1d	Male	2
12682	Work	Alone	Snowy	30	7AM	Bar	1d	Male	2
12683	Work	Alone	Sunny	80	7AM	Restaurant(20-50)	2h	Male	2

5 rows × 26 columns



In [6]: *#dimension of the data*
df.shape

Out[6]: (12684, 26)

- The data consists of 12684 rows and 26 columns.

In [7]: *# getting the column names of the data*
df.columns

Out[7]: Index(['destination', 'passanger', 'weather', 'temperature', 'time', 'coupon',
'expiration', 'gender', 'age', 'maritalStatus', 'has_children',
'education', 'occupation', 'income', 'car', 'Bar', 'CoffeeHouse',
'CarryAway', 'RestaurantLessThan20', 'Restaurant20To50',
'toCoupon_GEQ5min', 'toCoupon_GEQ15min', 'toCoupon_GEQ25min',
'direction_same', 'direction_opposite', 'Y'],
dtype='object')

```
In [8]: # summary of the data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12684 entries, 0 to 12683
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   destination                          12684 non-null  object
1   passanger                            12684 non-null  object
2   weather                              12684 non-null  object
3   temperature                          12684 non-null  int64
4   time                                 12684 non-null  object
5   coupon                              12684 non-null  object
6   expiration                          12684 non-null  object
7   gender                              12684 non-null  object
8   age                                 12684 non-null  object
9   maritalStatus                      12684 non-null  object
10  has_children                        12684 non-null  int64
11  education                          12684 non-null  object
12  occupation                          12684 non-null  object
13  income                             12684 non-null  object
14  car                                 108 non-null    object
15  Bar                                 12577 non-null  object
16  CoffeeHouse                        12467 non-null  object
17  CarryAway                          12533 non-null  object
18  RestaurantLessThan20               12554 non-null  object
19  Restaurant20To50                   12495 non-null  object
20  toCoupon_GEQ5min                   12684 non-null  int64
21  toCoupon_GEQ15min                  12684 non-null  int64
22  toCoupon_GEQ25min                  12684 non-null  int64
23  direction_same                      12684 non-null  int64
24  direction_opp                      12684 non-null  int64
25  Y                                   12684 non-null  int64
dtypes: int64(8), object(18)
memory usage: 2.5+ MB
```

- The DataFrame contains a total of 12,684 entries with 26 columns.
- Most columns have 12,684 non-null values, indicating they are complete.
- The majority of the data types are categorical (object), with a few numerical (int64) variables.
- Checked for unique values in each column in the dataframe.

```
In [9]: for col in df. columns:  
        print(df[col].value_counts(), '/n')
```

```
destination
No Urgent Place    6283
Home               3237
Work               3164
Name: count, dtype: int64 /n
passanger
Alone              7305
Friend(s)          3298
Partner            1075
Kid(s)             1006
Name: count, dtype: int64 /n
weather
Sunny              10069
Snowy              1405
Rainy              1210
Name: count, dtype: int64 /n
temperature
80                 6528
55                 3840
30                 2316
Name: count, dtype: int64 /n
time
6PM                3230
7AM                3164
10AM               2275
2PM                2009
10PM               2006
Name: count, dtype: int64 /n
coupon
Coffee House       3996
Restaurant(<20)     2786
Carry out & Take away 2393
Bar                2017
Restaurant(20-50)  1492
Name: count, dtype: int64 /n
expiration
1d                 7091
2h                 5593
Name: count, dtype: int64 /n
gender
Female             6511
Male               6173
Name: count, dtype: int64 /n
age
21                 2653
26                 2559
31                 2039
50plus             1788
36                 1319
41                 1093
46                 686
below21            547
Name: count, dtype: int64 /n
maritalStatus
Married partner    5100
Single             4752
Unmarried partner  2186
```

```

Divorced          516
Widowed           130
Name: count, dtype: int64 /n
has_children
0      7431
1      5253
Name: count, dtype: int64 /n
education
Some college - no degree          4351
Bachelors degree                 4335
Graduate degree (Masters or Doctorate) 1852
Associates degree                1153
High School Graduate             905
Some High School                 88
Name: count, dtype: int64 /n
occupation
Unemployed                      1870
Student                         1584
Computer & Mathematical         1408
Sales & Related                 1093
Education&Training&Library      943
Management                     838
Office & Administrative Support  639
Arts Design Entertainment Sports & Media 629
Business & Financial            544
Retired                        495
Food Preparation & Serving Related 298
Healthcare Practitioners & Technical 244
Healthcare Support              242
Community & Social Services      241
Legal                          219
Transportation & Material Moving 218
Architecture & Engineering       175
Personal Care & Service           175
Protective Service               175
Life Physical Social Science     170
Construction & Extraction        154
Installation Maintenance & Repair 133
Production Occupations           110
Building & Grounds Cleaning & Maintenance 44
Farming Fishing & Forestry       43
Name: count, dtype: int64 /n
income
$25000 - $37499      2013
$12500 - $24999      1831
$37500 - $49999      1805
$100000 or More      1736
$50000 - $62499      1659
Less than $12500      1042
$87500 - $99999       895
$75000 - $87499       857
$62500 - $74999       846
Name: count, dtype: int64 /n
car
Scooter and motorcycle    22
Mazda5                    22
do not drive              22

```

```

crossover 21
Car that is too old to install Onstar :D 21
Name: count, dtype: int64 /n
Bar
never 5197
less1 3482
1~3 2473
4~8 1076
gt8 349
Name: count, dtype: int64 /n
CoffeeHouse
less1 3385
1~3 3225
never 2962
4~8 1784
gt8 1111
Name: count, dtype: int64 /n
CarryAway
1~3 4672
4~8 4258
less1 1856
gt8 1594
never 153
Name: count, dtype: int64 /n
RestaurantLessThan20
1~3 5376
4~8 3580
less1 2093
gt8 1285
never 220
Name: count, dtype: int64 /n
Restaurant20To50
less1 6077
1~3 3290
never 2136
4~8 728
gt8 264
Name: count, dtype: int64 /n
toCoupon_GEQ5min
1 12684
Name: count, dtype: int64 /n
toCoupon_GEQ15min
1 7122
0 5562
Name: count, dtype: int64 /n
toCoupon_GEQ25min
0 11173
1 1511
Name: count, dtype: int64 /n
direction_same
0 9960
1 2724
Name: count, dtype: int64 /n
direction_opp
1 9960
0 2724
Name: count, dtype: int64 /n

```



```

Y
1    7210
0    5474
Name: count, dtype: int64 /n

```

- All the columns contain categorical values with majority having the object datatype and few numerical columns.

```

In [10]: # checking the percentage of missing values
df.isnull().sum()/len(df)*100

```

```

Out[10]:
0

```

destination	0.000000
passanger	0.000000
weather	0.000000
temperature	0.000000
time	0.000000
coupon	0.000000
expiration	0.000000
gender	0.000000
age	0.000000
maritalStatus	0.000000
has_children	0.000000
education	0.000000
occupation	0.000000
income	0.000000
car	99.148534
Bar	0.843582
CoffeeHouse	1.710817
CarryAway	1.190476
RestaurantLessThan20	1.024913
Restaurant20To50	1.490066
toCoupon_GEQ5min	0.000000
toCoupon_GEQ15min	0.000000
toCoupon_GEQ25min	0.000000
direction_same	0.000000
direction_opp	0.000000
Y	0.000000

dtype: float64

- Most of the columns do not have missing values. However the 'car' column contains a huge percentage of missing values. This column was dropped because imputing values would increase the risk of inaccuracies.
- Other columns that contain null values are; Bar, CoffeeHouse, carryAway, RestaurantLessThan20 and Restaurant20To50. Since these columns had a small percentage of missing values, they were handled through imputation.

```
In [11]: # dropping the car column
df.drop('car', axis=1, inplace = True)
```

```
In [12]: #list of columns with null values
null_cols = ['Bar', 'CoffeeHouse', 'CarryAway', 'RestaurantLessThan20', 'Restau
for col in null_cols:
    print(col, df[col].value_counts()) #count of unique entries in each column
```

```
Bar Bar
never      5197
less1      3482
1~3        2473
4~8        1076
gt8         349
Name: count, dtype: int64
CoffeeHouse CoffeeHouse
less1      3385
1~3        3225
never      2962
4~8        1784
gt8        1111
Name: count, dtype: int64
CarryAway CarryAway
1~3        4672
4~8        4258
less1      1856
gt8        1594
never       153
Name: count, dtype: int64
RestaurantLessThan20 RestaurantLessThan20
1~3        5376
4~8        3580
less1      2093
gt8        1285
never       220
Name: count, dtype: int64
Restaurant20To50 Restaurant20To50
less1      6077
1~3        3290
never      2136
4~8        728
gt8        264
Name: count, dtype: int64
```

- Since all the columns with missing values are categorical, the missing values were imputed using the mode.

```
In [13]: #replacing null values in the categorical columns with the mode
for col in null_cols:
    df[col].fillna(df[col].mode()[0], inplace=True)
```

- Checked for unique values in all the categorical columns to confirm whether some cleaning was necessary.

```
In [14]: for col in df.select_dtypes(include='object').columns:
        print(f"Unique values in column '{col}': {df[col].unique()}") #getting uni

Unique values in column 'destination': ['No Urgent Place' 'Home' 'Work']
Unique values in column 'passanger': ['Alone' 'Friend(s)' 'Kid(s)' 'Partner']
Unique values in column 'weather': ['Sunny' 'Rainy' 'Snowy']
Unique values in column 'time': ['2PM' '10AM' '6PM' '7AM' '10PM']
Unique values in column 'coupon': ['Restaurant(<20)' 'Coffee House' 'Carry ou
t & Take away' 'Bar'
'Restaurant(20-50)']
Unique values in column 'expiration': ['1d' '2h']
Unique values in column 'gender': ['Female' 'Male']
Unique values in column 'age': ['21' '46' '26' '31' '41' '50plus' '36' 'below
21']
Unique values in column 'maritalStatus': ['Unmarried partner' 'Single' 'Marri
ed partner' 'Divorced' 'Widowed']
Unique values in column 'education': ['Some college - no degree' 'Bachelors d
egree' 'Associates degree'
'High School Graduate' 'Graduate degree (Masters or Doctorate)'
'Some High School']
Unique values in column 'occupation': ['Unemployed' 'Architecture & Engineeri
ng' 'Student'
'Education&Training&Library' 'Healthcare Support'
'Healthcare Practitioners & Technical' 'Sales & Related' 'Management'
'Arts Design Entertainment Sports & Media' 'Computer & Mathematical'
'Life Physical Social Science' 'Personal Care & Service'
'Community & Social Services' 'Office & Administrative Support'
'Construction & Extraction' 'Legal' 'Retired'
'Installation Maintenance & Repair' 'Transportation & Material Moving'
'Business & Financial' 'Protective Service'
'Food Preparation & Serving Related' 'Production Occupations'
'Building & Grounds Cleaning & Maintenance' 'Farming Fishing & Forestry']
Unique values in column 'income': ['$37500 - $49999' '$62500 - $74999' '$1250
0 - $24999' '$75000 - $87499'
'$50000 - $62499' '$25000 - $37499' '$100000 or More' '$87500 - $99999'
'Less than $12500']
Unique values in column 'Bar': ['never' 'less1' '1~3' 'gt8' '4~8']
Unique values in column 'CoffeeHouse': ['never' 'less1' '4~8' '1~3' 'gt8']
Unique values in column 'CarryAway': ['1~3' '4~8' 'gt8' 'less1' 'never']
Unique values in column 'RestaurantLessThan20': ['4~8' '1~3' 'less1' 'gt8' 'n
ever']
Unique values in column 'Restaurant20To50': ['1~3' 'less1' 'never' 'gt8' '4~
8']
```


- For uniformity, better understanding and reduction of bulkiness of the values the following values in their respective columns were cleaned.

```
In [15]: #cleaning the columns
df['age'] = df['age'].replace({'50plus': '50+',
                              'below21': '<21'})
df['maritalStatus'] = df['maritalStatus'].replace({
    'Unmarried partner': 'Unmarried',
    'Married partner': 'Married'
})
df['income'] = df['income'].replace({
    'Less than $12500': '< $12500',
    '$100000 or More': '$100,000+'
})
for col in ['Bar', 'CoffeeHouse', 'CarryAway', 'RestaurantLessThan20', 'Restau
df[col] = df[col].replace({
    'less1': '<1',      # Replace 'less1' with '<1'
    '1~3': '1-3',      # Replace '1~3' with '1-3'
    '4~8': '4-8',      # Replace '4~8' with '4-8'
    'gt8': '>8'         # Replace 'gt8' with '>8'
})
df.head()
```

```
Out[15]:
```

	destination	passanger	weather	temperature	time	coupon	expiration	gender	age
0	No Urgent Place	Alone	Sunny	55	2PM	Restaurant(<20)	1d	Female	21
1	No Urgent Place	Friend(s)	Sunny	80	10AM	Coffee House	2h	Female	21
2	No Urgent Place	Friend(s)	Sunny	80	10AM	Carry out & Take away	2h	Female	21
3	No Urgent Place	Friend(s)	Sunny	80	2PM	Coffee House	2h	Female	21
4	No Urgent Place	Friend(s)	Sunny	80	2PM	Coffee House	1d	Female	21

5 rows × 25 columns



- Checked of unique values in the numerical columns. The entries were uniform so no further cleaning was needed.

```
In [16]: for col in df.select_dtypes(include='number').columns:  
         print(f"Unique values in column '{col}': {df[col].unique()}")
```

```
Unique values in column 'temperature': [55 80 30]  
Unique values in column 'has_children': [1 0]  
Unique values in column 'toCoupon_GEQ5min': [1]  
Unique values in column 'toCoupon_GEQ15min': [0 1]  
Unique values in column 'toCoupon_GEQ25min': [0 1]  
Unique values in column 'direction_same': [0 1]  
Unique values in column 'direction_opp': [1 0]  
Unique values in column 'Y': [1 0]
```

- Explored the temperature column further by viewing its statistics.

```
In [17]: df.temperature.describe().T
```

```
Out[17]:
```

	temperature
count	12684.000000
mean	63.301798
std	19.154486
min	30.000000
25%	55.000000
50%	80.000000
75%	80.000000
max	80.000000

dtype: float64

- Discovered that there were some events where the temperature was 30 farhreneit and the weather was sunny at the same time. 30 farhreneit is equivalent to -1 degree celcius.
- From further research it is possible for the temperature to be -1°C (30.2°F) and sunny at the same time. Low temperatures don't prevent the sun from shining. SO the entries will remain as they are.


```
In [18]: # get the data where temp is 30 and weather is sunny
# Filter the DataFrame
filtered_df = df[(df['temperature'] == 30) & (df['weather'] == 'Sunny')]

# len of filtered_df
filtered_df
```

```
Out[18]:
```

	destination	passanger	weather	temperature	time	coupon	expiration	gender
6586	No Urgent Place	Friend(s)	Sunny	30	10PM	Coffee House	1d	Male
6601	Work	Alone	Sunny	30	7AM	Bar	1d	Male
6608	No Urgent Place	Friend(s)	Sunny	30	10PM	Coffee House	1d	Male
6623	Work	Alone	Sunny	30	7AM	Bar	1d	Male
6628	No Urgent Place	Friend(s)	Sunny	30	10PM	Coffee House	1d	Male
...
12665	No Urgent Place	Friend(s)	Sunny	30	10AM	Carry out & Take away	2h	Male
12669	No Urgent Place	Partner	Sunny	30	10AM	Restaurant(20-50)	1d	Male
12673	Home	Alone	Sunny	30	6PM	Carry out & Take away	1d	Male
12677	Home	Partner	Sunny	30	6PM	Restaurant(<20)	1d	Male
12678	Home	Partner	Sunny	30	10PM	Restaurant(<20)	2h	Male

911 rows × 25 columns

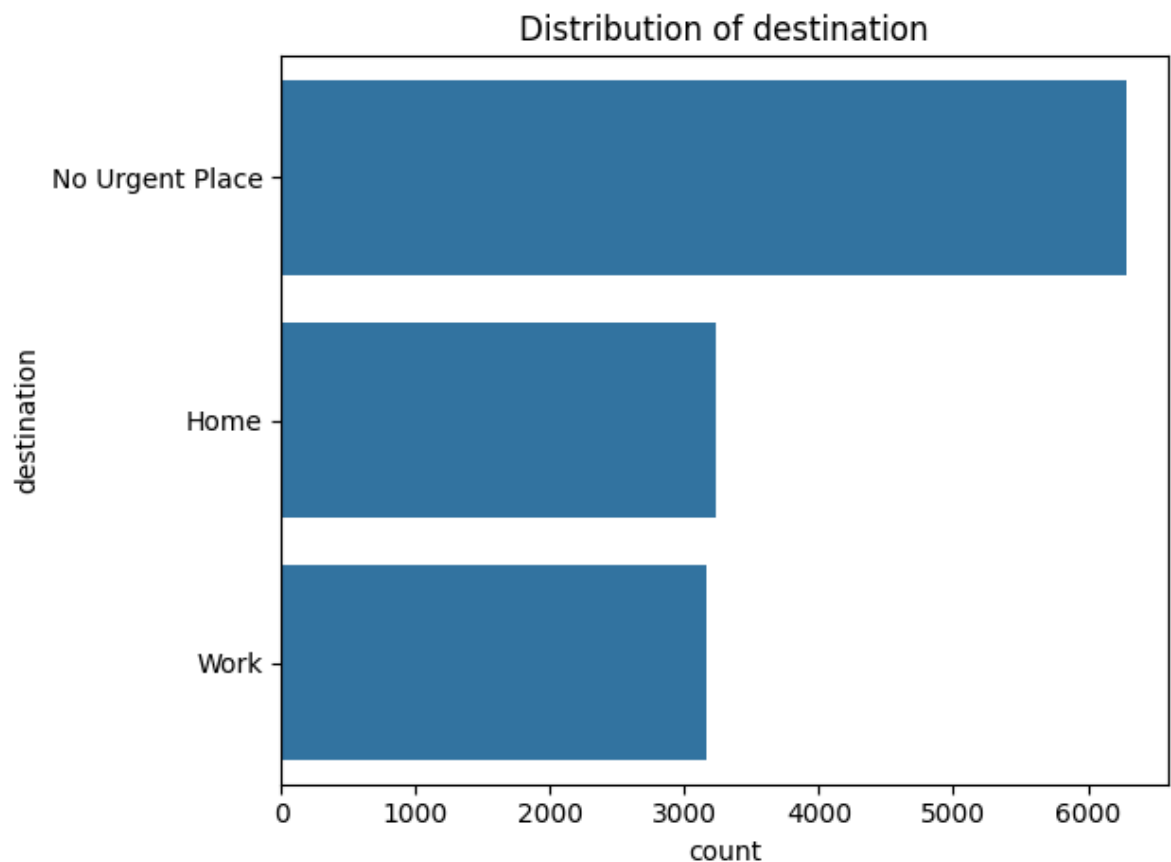




EDA

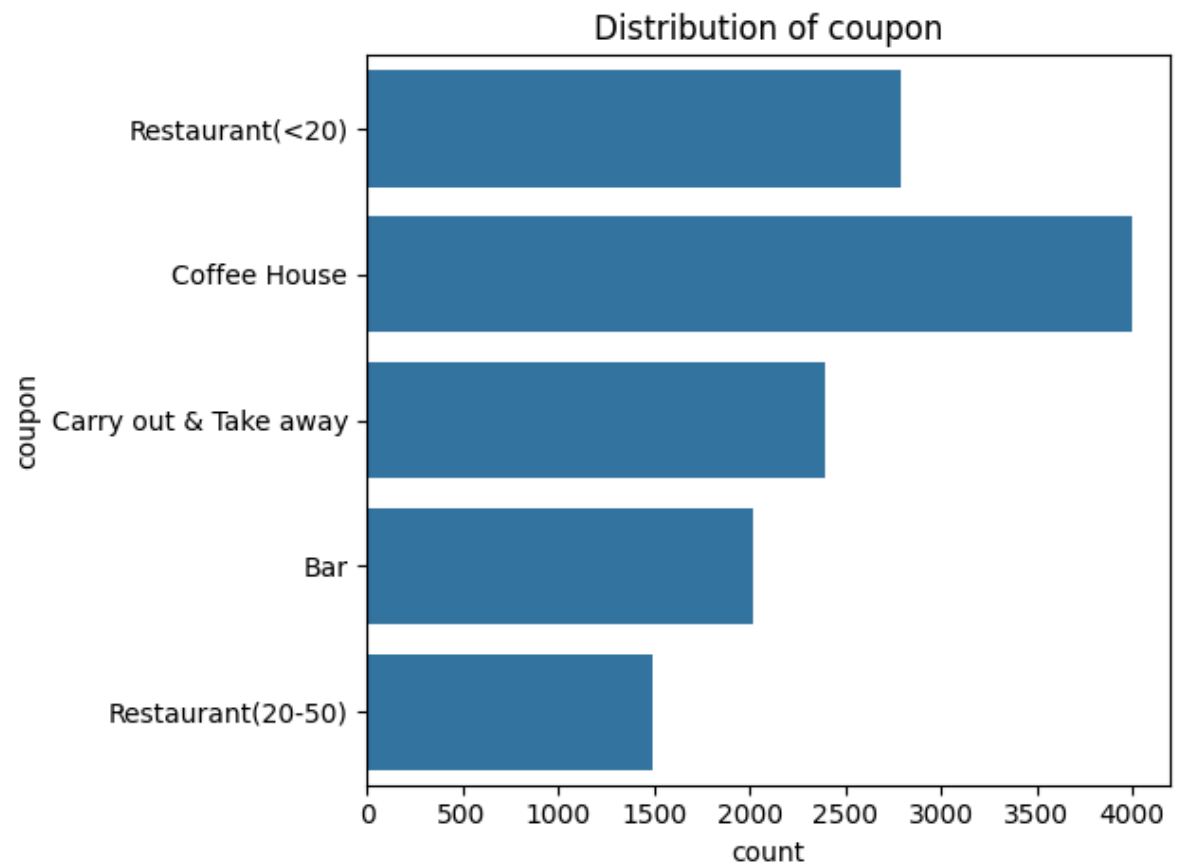

Univariate Analysis



```
In [ ]: for col in df.select_dtypes(include='object').columns:
        sns.countplot(df[col])
        plt.title(f'Distribution of {col}')
        plt.tight_layout() # Adjust layout to prevent labels from overlapping
        plt.show()

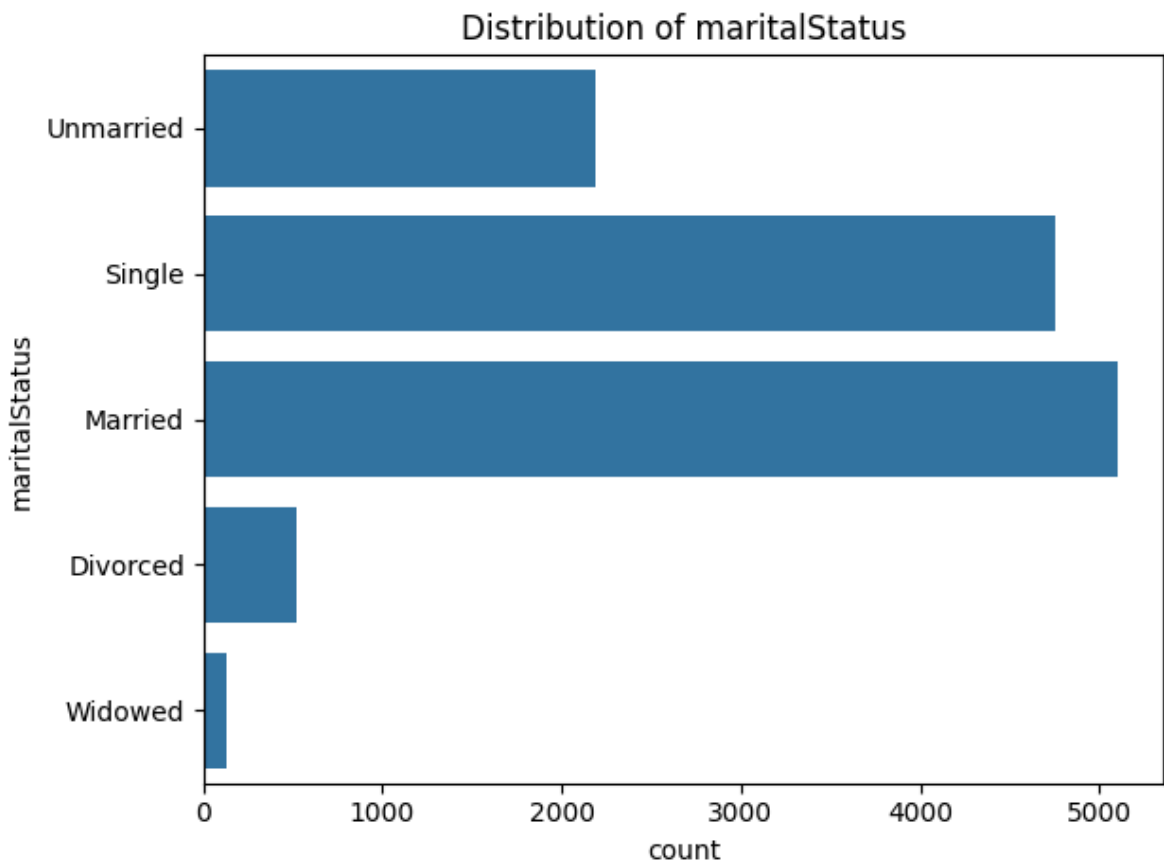

plt.tight_layout() # Adjust layout to prevent overlapping
plt.show()
```


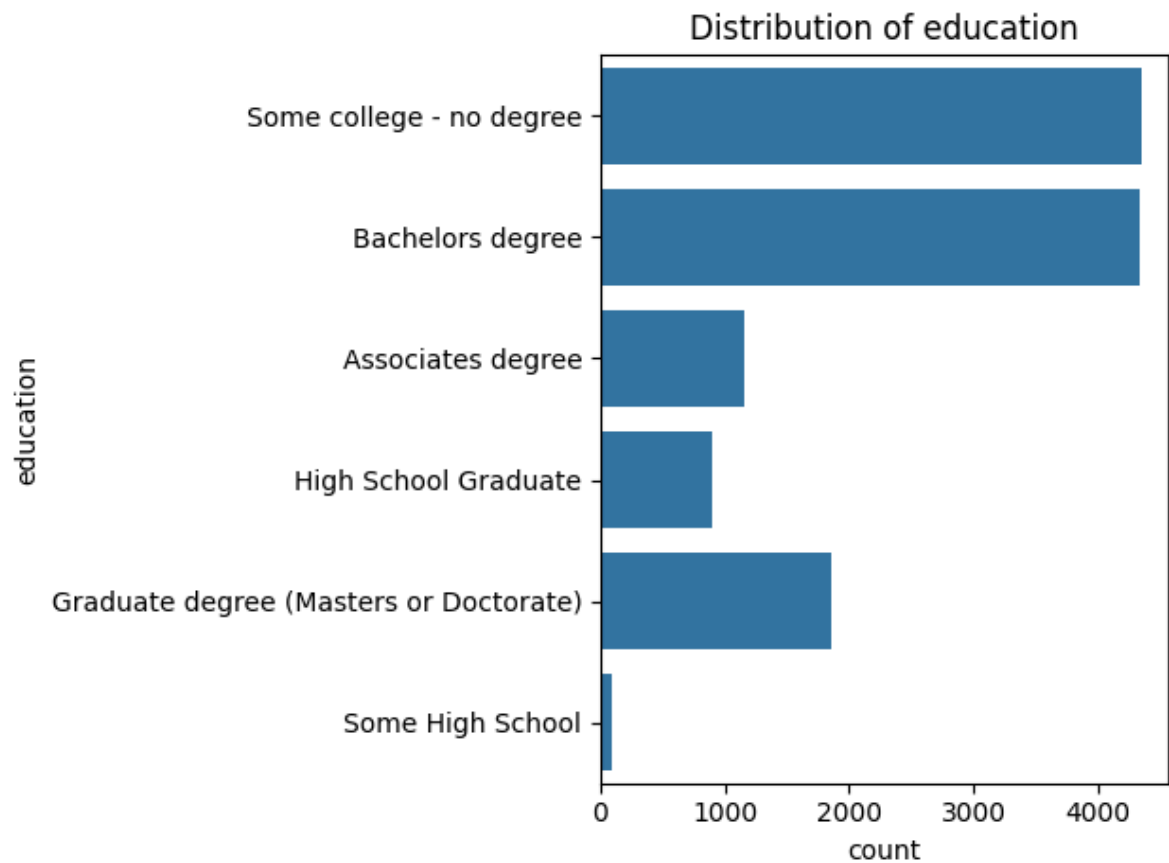


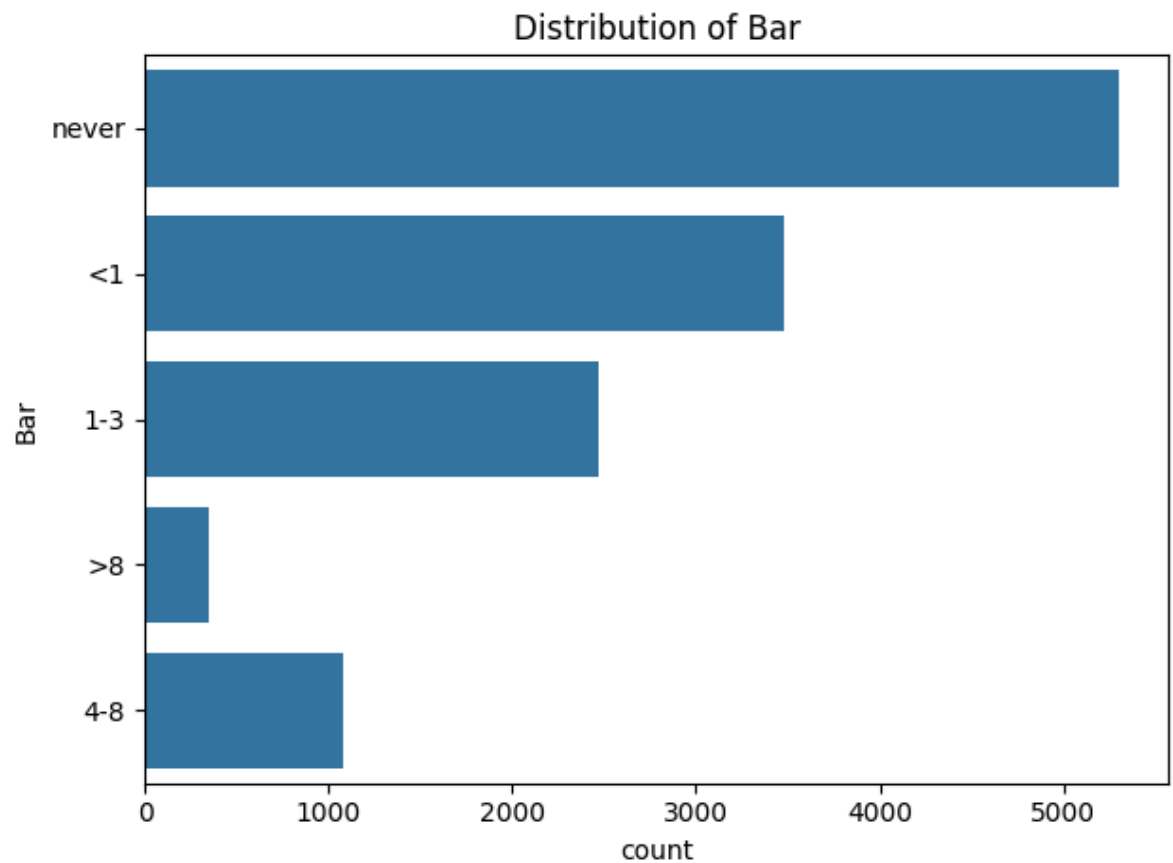



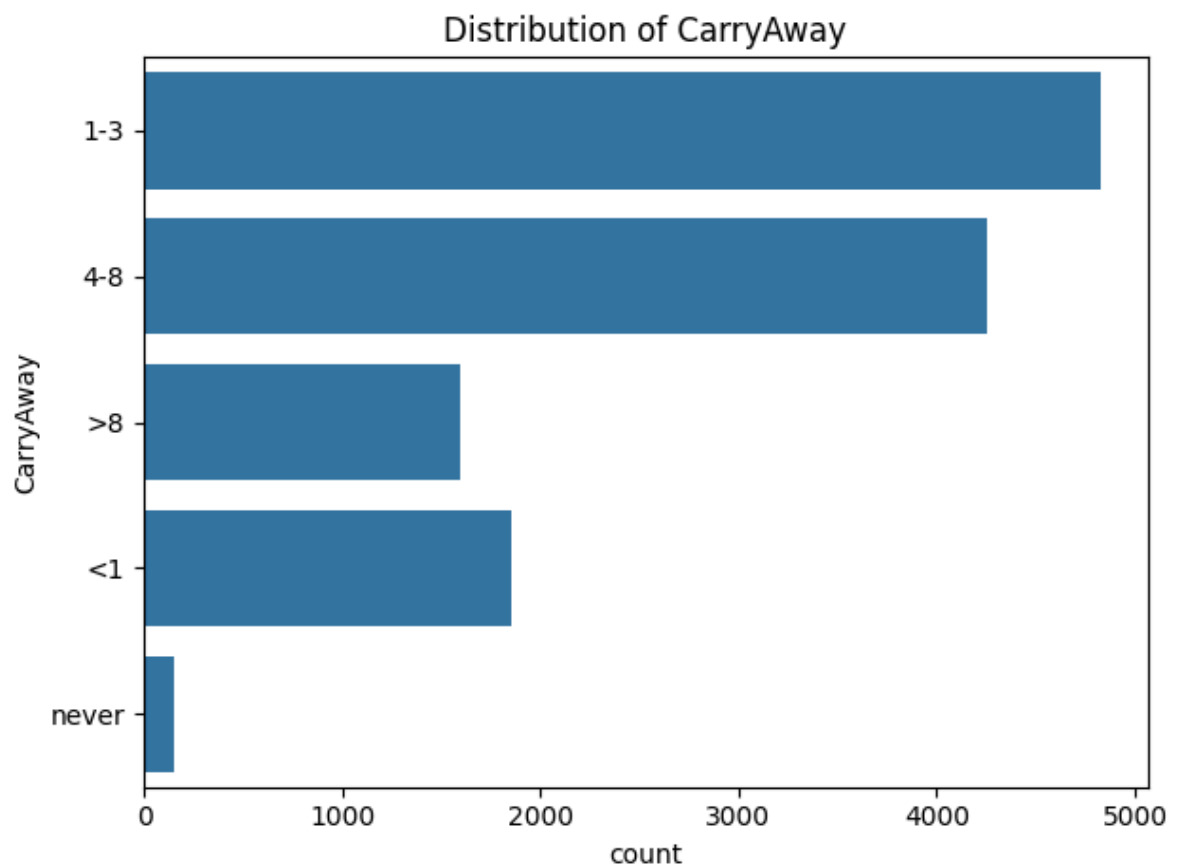
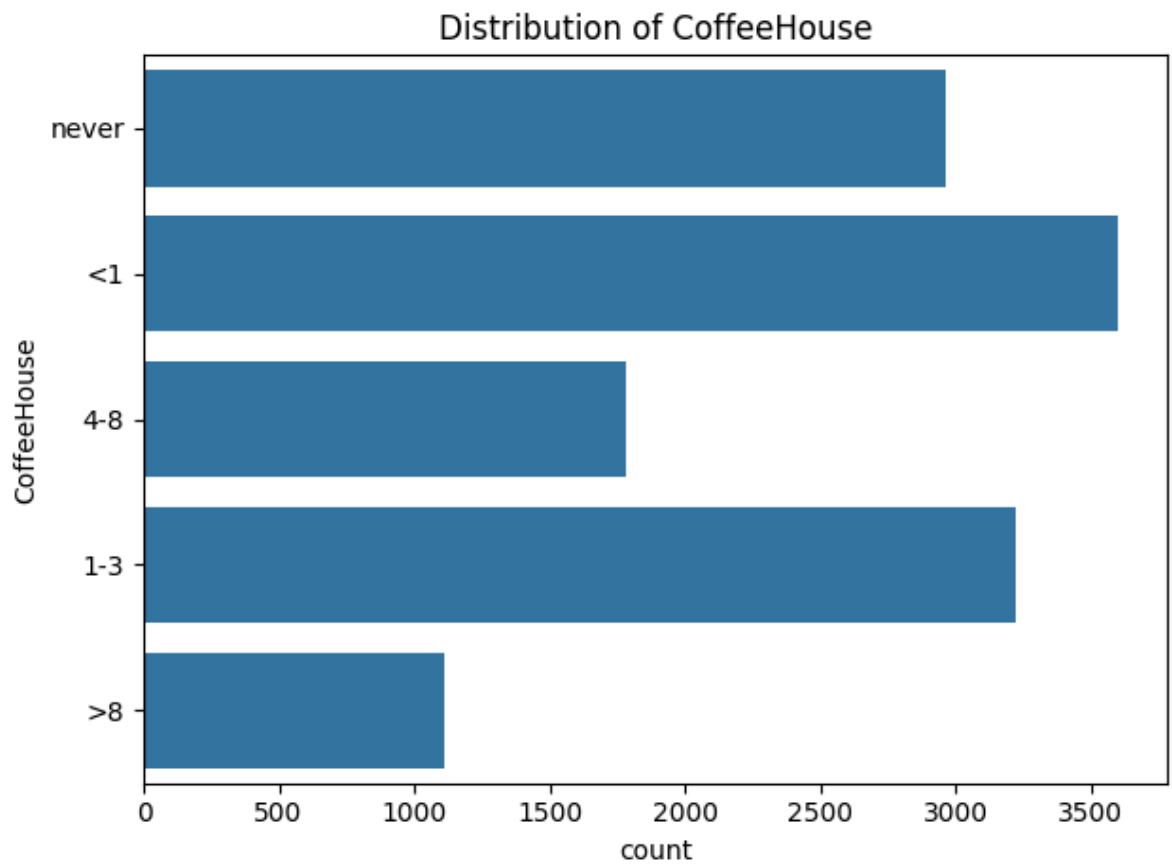


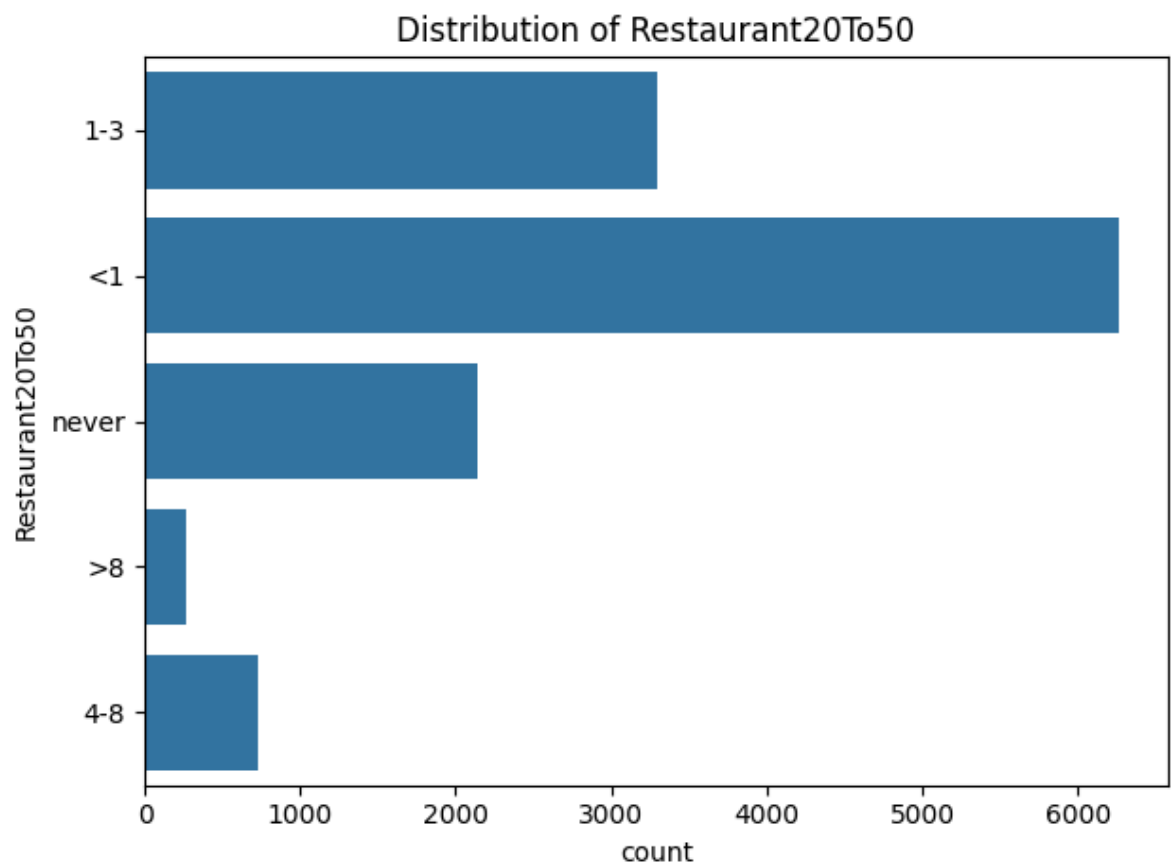
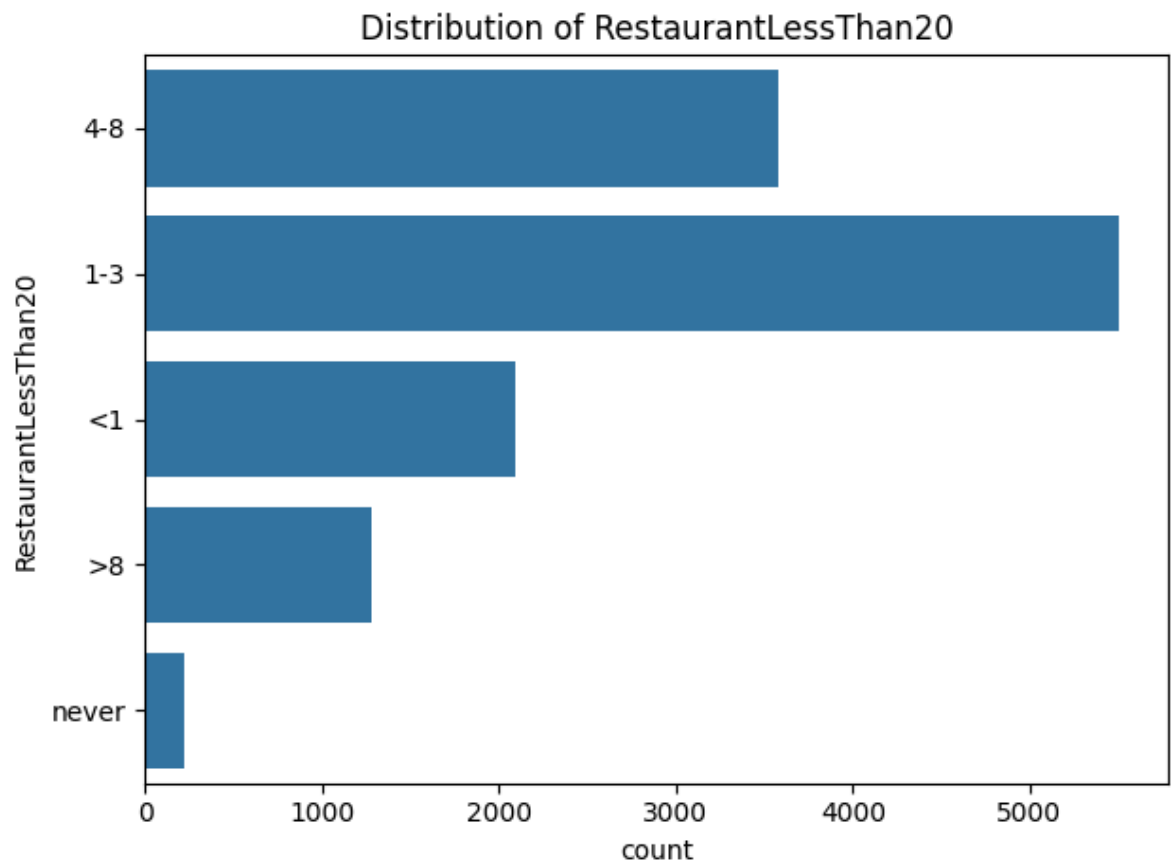












<Figure size 640x480 with 0 Axes>

Observations

- Most passengers that boarded had no urgent place to go.
- A high percentage of passengers were travelling alone.
- The weather was sunny during most of the travel time.
- Most of the passengers were between the age group of 20-30 years.
- The marital status of most passengers was 'Married'.
- The income of the passengers varied over a wide range. There was no significant variation in the passengers' income that were on board.
- A high percentage of passengers had never visited the bar.
- Most passengers appear to have a moderate coffee consumption level,

```
In [ ]: # distribution of the y in relation to destination
# Define the columns to plot
cols = ['destination', 'age', 'temperature', 'gender', 'coupon', 'expiration']

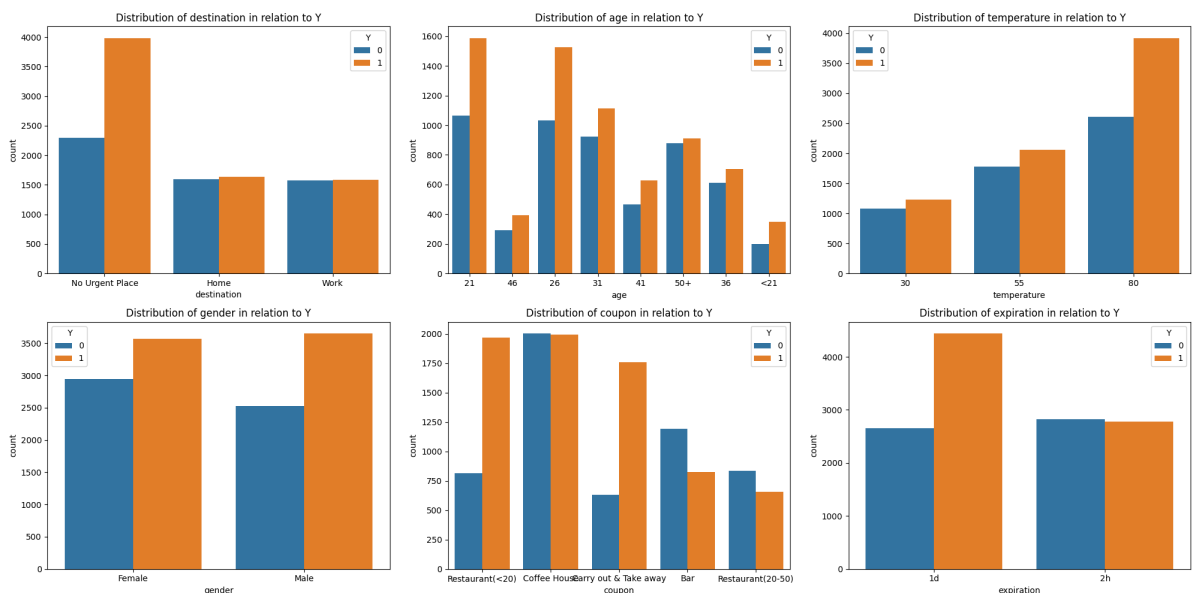
# Set the number of columns for the grid
n_cols = 3
n_rows = (len(cols) + n_cols - 1) // n_cols # Calculate the number of rows needed

# Create subplots
fig, axes = plt.subplots(n_rows, n_cols, figsize=(20, n_rows * 5))
axes = axes.flatten() # Flatten to easily iterate

# Loop through columns and plot
for i, col in enumerate(cols):
    sns.countplot(x=df[col], hue='Y', data=df, ax=axes[i])
    axes[i].set_title(f'Distribution of {col} in relation to Y')

# Hide any unused axes
for j in range(i + 1, n_rows * n_cols):
    axes[j].axis('off')

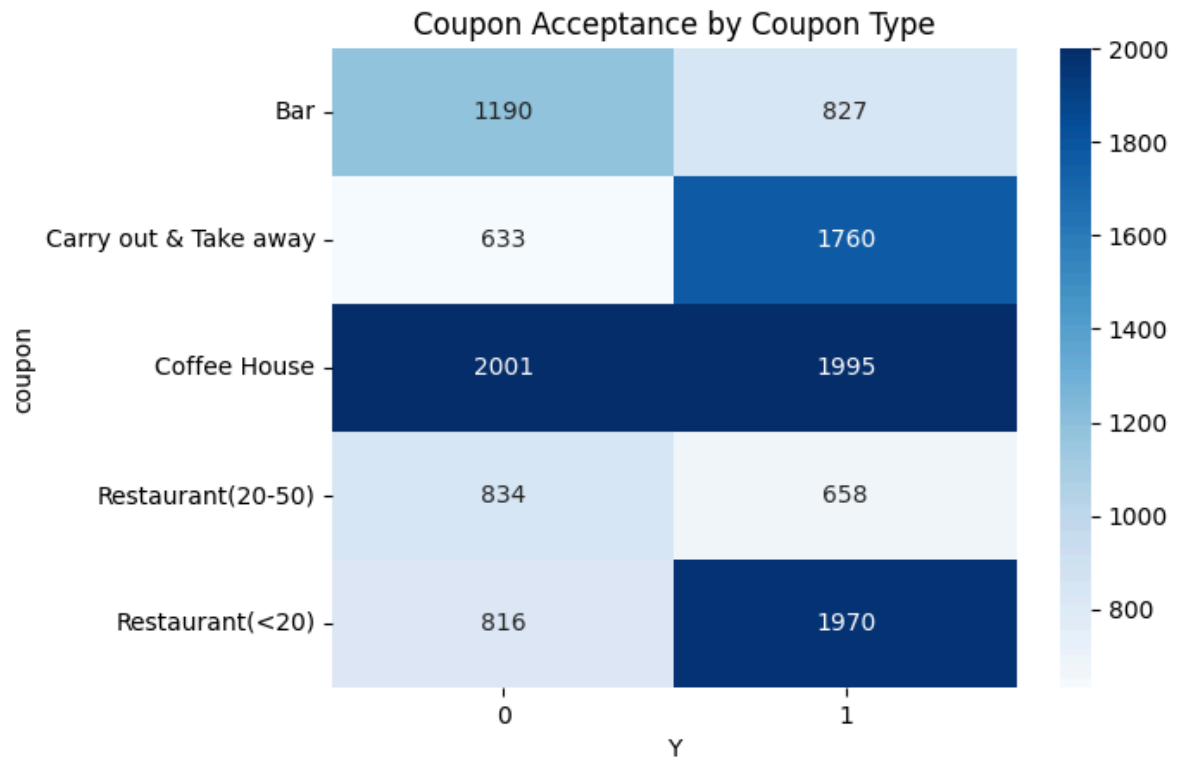
plt.tight_layout()
plt.show()
```




Observations

- Passengers traveling to "No Urgent Place" are the most frequent, indicating a higher likelihood of coupon acceptance in scenarios where they have no specific destination. This may suggest a tendency to explore options when they are not in a hurry. The counts for "Home" and "Work" are lower, but there is a noticeable difference in coupon acceptance between the two. Passengers going "Home" seem to be more likely to accept coupons than those heading to "Work," indicating that after a workday, passengers may be more inclined to consider promotional offers.
- Younger passengers (especially those in the <21 age group) appear to have lower acceptance rates for coupons compared to those in the 21-31 age groups. This may suggest that older, more established individuals are more responsive to coupon offers.
- The trend indicates that as the temperature increases, the number of coupon acceptances (1) also increases, suggesting that passengers are more likely to accept coupons in warmer weather. This could imply that people are more likely to go out and dine or shop when it's warmer, thus being more receptive to coupons.
- Males show a slightly higher acceptance rate for coupons compared to females. This might suggest that male passengers are more responsive to promotions or incentives than their female counterparts. Overall, the distribution remains fairly balanced, indicating a similar inclination towards accepting or rejecting coupons across genders.
- The "Restaurant(<20)" coupon shows the highest acceptance count, indicating its effectiveness in enticing passengers. Other coupon types like "Carry out & Take away" and "coffeeHouse" also attract a notable number of acceptances, suggesting that price-sensitive offers resonate well with customers. The "Restaurant(20 - 50)" and "Bar" categories have more rejections, indicating that these offers might not appeal to as many passengers.
- Coupons with a longer expiration time (24 hours/1d) have a higher count of accepted coupons.

```
In [ ]: coupon_acceptance = pd.crosstab(df['coupon'], df['Y'])  
sns.heatmap(coupon_acceptance, annot=True, cmap='Blues', fmt='d')  
plt.title('Coupon Acceptance by Coupon Type')  
plt.show()
```



```
In [ ]: sns.heatmap(df.select_dtypes(include='number').corr(), annot=True, cmap='Blues')
plt.show()
```



Data Preprocessing

- The first step was to define our target variable which was column 'Y'. This column has two entries; 0 and 1. 0 implies that the passenger didn't receive a coupon while 1 implies that the passenger received a coupon.
- All the remaining columns were assigned to the variable X.
- Split the data into training sets and test sets.
- Preprocessing steps were done separately on the train and test sets to avoid data leakage.
- Transformed all the columns to numerical format so as to run the ml algorithm models efficiently.
- Scaled the train and test sets so that all features contribute equally to the ml models.

```
In [19]: # define X and y variables
X = df.drop('Y', axis=1)
y = df['Y']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# One hot encoding the train and test data separately using the sklearn library
ohe = OneHotEncoder(sparse_output=False) # set sparse_output to False
X_train = ohe.fit_transform(X_train)
X_test = ohe.transform(X_test)
# transform the X_train and X_test to dataframe
X_train = pd.DataFrame(X_train, columns=ohe.get_feature_names_out(X.columns))
X_test = pd.DataFrame(X_test, columns=ohe.get_feature_names_out(X.columns))
```

```
In [20]: #scaling the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

- Saved the new datasets used for modelling in csv files.

```
In [21]: # transform the X_train and X_test to dataframe
X_train = pd.DataFrame(X_train, columns=ohe.get_feature_names_out(X.columns))
X_test = pd.DataFrame(X_test, columns=ohe.get_feature_names_out(X.columns))
X_train.to_csv('cleaned_traindata.csv', index=False)
X_test.to_csv('cleaned_testdata.csv', index=False)
```

Modelling

- Logistic Regression was used as the baseline model due to its simplicity, ease of interpretation, it also serves as a reliable starting point for evaluating model performance.
- Other ML algorithms used were Decision Trees and RandomForestClassifier.
- For each model, a classification report was generated to show how the model was performing based on all the metrics of evaluation used for these algorithms.

Metric of Evaluation

- **Recall** :Our goal is to maximise the coupon usage ie.to maximize true positives (correctly identifying customers who will use a coupon)

Logistic Regression

```
In [22]: # Baseline model
lr = LogisticRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_train)
print("Train Classification Report\n",classification_report(y_train, y_pred))
y_predr1 = lr.predict(X_test)
print("Test Classification Report\n",classification_report(y_test, y_predr1))
```

Train Classification Report

	precision	recall	f1-score	support
0	0.66	0.58	0.62	4346
1	0.71	0.78	0.74	5801
accuracy			0.69	10147
macro avg	0.69	0.68	0.68	10147
weighted avg	0.69	0.69	0.69	10147

Test Classification Report

	precision	recall	f1-score	support
0	0.67	0.57	0.62	1128
1	0.69	0.78	0.73	1409
accuracy			0.69	2537
macro avg	0.68	0.67	0.68	2537
weighted avg	0.68	0.69	0.68	2537

Recall for Class 0 (Non-Users of the Coupon):

Training Set: 0.58 Test Set: 0.57

This recall score indicates that the model is correctly identifying about 57-58% of actual non-users as non-users. In other words, there's a significant portion of non-users (false negatives) that the model incorrectly predicts as coupon users.

Recall for Class 1 (Users of the Coupon):

Training Set: 0.78 Test Set: 0.78

The model correctly identifies about 78% of actual coupon users as users, which is a relatively good recall. This means the model is better at identifying those who will use the coupon, leading to fewer missed opportunities (false negatives) among coupon users.

Decision Trees

```
In [ ]: #Decision Trees
tree = DecisionTreeClassifier()
tree.fit(X_train, y_train)
y_pred2 = tree.predict(X_train)
print("Train Classification Report\n",classification_report(y_train, y_pred2))
y_pred3 = tree.predict(X_test)
print("Test Classification Report\n",classification_report(y_test, y_pred3))
```

Train Classification Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	4346
1	1.00	1.00	1.00	5801
accuracy			1.00	10147
macro avg	1.00	1.00	1.00	10147
weighted avg	1.00	1.00	1.00	10147

Test Classification Report

	precision	recall	f1-score	support
0	0.65	0.63	0.64	1128
1	0.71	0.73	0.72	1409
accuracy			0.69	2537
macro avg	0.68	0.68	0.68	2537
weighted avg	0.68	0.69	0.68	2537

- The model achieved a precision, recall, and F1-score of 1.00 for both classes on the training set. This perfect accuracy (1.00) indicates the model memorized the training data, fitting it exactly.
- On the test set, the accuracy dropped to 0.69, with similar drops in precision, recall, and F1-score.
- The drastic difference between the training and test set metrics highlights that the decision tree model is highly overfit to the training data, struggling to generalize to unseen data.

RandomForestClassifier

```
In [ ]: # RandomForestClassifier
rf= RandomForestClassifier()
rf.fit(X_train, y_train)
y_pred4 = rf.predict(X_train)
print("Train Classification Report\n",classification_report(y_train, y_pred4))
y_pred5 = rf.predict(X_test)
print("Train Classification Report\n",classification_report(y_test, y_pred5))
```

Train Classification Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	4346
1	1.00	1.00	1.00	5801
accuracy			1.00	10147
macro avg	1.00	1.00	1.00	10147
weighted avg	1.00	1.00	1.00	10147

Train Classification Report

	precision	recall	f1-score	support
0	0.74	0.64	0.69	1128
1	0.74	0.82	0.78	1409
accuracy			0.74	2537
macro avg	0.74	0.73	0.73	2537
weighted avg	0.74	0.74	0.74	2537

- While the test set performance is not as high as the training set, it suggests the model is generalizing better than a single decision tree. However, the difference in scores between the training and test sets still suggests overfitting.

Fine tuning the models

- Fine-tuned the models to improve their performance by optimizing hyperparameters, which can reduce overfitting or underfitting, and ultimately produce more reliable results on new, unseen data

```

In [ ]: #Finetuning the models
# Parameter grid for GridSearchCV
param_grid = {
    'C': [0.00001, 0.0001, 0.001, 0.01, 0.1, 1], # Regularization strength
    'penalty': ['l1', 'l2', 'elasticnet'], # Penalty types
    'solver': ['liblinear', 'saga'] # Solver for optimization
}
grid = GridSearchCV(lr, param_grid, cv=5, scoring= 'precision')
grid.fit(X_train, y_train)
print("Best parameters for Logistic Regression:", grid.best_params_)
# Predict on the test set using the best model
best_logreg = grid.best_estimator_
y_predt = best_logreg.predict(X_test)
y_pred0 = best_logreg.predict(X_train)

# classification report for both train and test data
print("Train Classification Report\n", classification_report(y_train, y_pred0))
print("Test Classification Report\n", classification_report(y_test, y_predt))

```

Best parameters for Logistic Regression: {'C': 0.001, 'penalty': 'l2', 'solver': 'liblinear'}

Train Classification Report

	precision	recall	f1-score	support
0	0.65	0.59	0.62	4346
1	0.71	0.76	0.74	5801
accuracy			0.69	10147
macro avg	0.68	0.68	0.68	10147
weighted avg	0.69	0.69	0.69	10147

Test Classification Report

	precision	recall	f1-score	support
0	0.66	0.57	0.62	1128
1	0.69	0.77	0.73	1409
accuracy			0.68	2537
macro avg	0.68	0.67	0.67	2537
weighted avg	0.68	0.68	0.68	2537

- Good Recall for Users: The model has a high recall for class 1 (coupon users), capturing about 77% of actual coupon users in the test set, which can help increase coupon redemption rates.
- Moderate Recall for Non-Users: The model's recall for non-users is moderate, so it may misclassify some non-users as users.
- There was no much improvement from the baseline model .
- Note that the model only perform as good as the parameters fed into the dictionary grid.


```

In [ ]: #random forest using gridsearchcv
param_grid = {'n_estimators':[100, 10],
              'max_depth': [10, 20, 30],
              'min_samples_split': [10, 20],
              'min_samples_leaf': [10, 20],
              'max_features': ['sqrt', 'log2', None],
              'bootstrap': [True, False],
              'class_weight': ['balanced', 'balanced_subsample', None],
              'random_state': [42]
             }

grid = GridSearchCV(rf, param_grid, cv=5, scoring= 'precision')
grid.fit(X_train, y_train)
print("Best parameters for Random Forest:", grid.best_params_)

best_rf = grid.best_estimator_
y_pred = best_rf.predict(X_test)
y_pred1 = best_rf.predict(X_train)
print(classification_report(y_train, y_pred1))
print(classification_report(y_test, y_pred))

```

Best parameters for Random Forest: {'bootstrap': False, 'class_weight': 'balanced', 'max_depth': 30, 'max_features': 'sqrt', 'min_samples_leaf': 10, 'min_samples_split': 10, 'n_estimators': 100, 'random_state': 42}

	precision	recall	f1-score	support
0	0.78	0.81	0.79	4346
1	0.85	0.83	0.84	5801
accuracy			0.82	10147
macro avg	0.81	0.82	0.81	10147
weighted avg	0.82	0.82	0.82	10147

	precision	recall	f1-score	support
0	0.68	0.69	0.69	1128
1	0.75	0.74	0.75	1409
accuracy			0.72	2537
macro avg	0.72	0.72	0.72	2537
weighted avg	0.72	0.72	0.72	2537

- The model shows great improvement from the previous RandomForestClassifier model which was highly overfitting.
- Although there is still a huge gap between the recall metrics of the train and test sets which implies that the model could still be overfitting.

```

In [ ]: # Define the parameter grid
# DecisionTree with gridsearchcv
param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10, 20],
    'min_samples_leaf': [1, 2, 4, 6, 8]
}
grid = GridSearchCV(tree, param_grid, cv=5, scoring= 'precision')
grid.fit(X_train, y_train)
print("Best parameters for Decision Tree:", grid.best_params_)
#classification report for train and test data
best_tree = grid.best_estimator_
y_pred2 = best_tree.predict(X_train)
y_pred3 = best_tree.predict(X_test)

print('Train classificatio Report',classification_report(y_train, y_pred2))
print('Test classification Report',classification_report(y_test, y_pred3))

```

Best parameters for Decision Tree: {'criterion': 'gini', 'max_depth': 30, 'min_samples_leaf': 2, 'min_samples_split': 5}

Train classificatio Report	precision	recall	f1-score	support
----------------------------	-----------	--------	----------	---------

0	0.87	0.96	0.91	4346
1	0.97	0.89	0.93	5801

accuracy			0.92	10147
macro avg	0.92	0.93	0.92	10147
weighted avg	0.93	0.92	0.92	10147

Test classification Report	precision	recall	f1-score	support
----------------------------	-----------	--------	----------	---------

0	0.64	0.69	0.66	1128
1	0.73	0.68	0.71	1409

accuracy			0.69	2537
macro avg	0.68	0.69	0.68	2537
weighted avg	0.69	0.69	0.69	2537


- The model did not show much great improvement from the hyperparameter tuning. The Model is still highly overfitting, which suggests that the model is learning the training data too well with minimal generalisation.

```

In [ ]: # plotting the AUC and ROC curve
#compute the ROC curve and roc_auc score
fpr, tpr, thresholds = roc_curve(y_test, y_predr1)
roc_auc = roc_auc_score(y_test, y_predr1)
fpr2, tpr2, thresholds2 = roc_curve(y_test, y_pred)
roc_auc2 = roc_auc_score(y_test, y_pred)
fpr3, tpr3, thresholds3 = roc_curve(y_test, y_pred3)
roc_auc3 = roc_auc_score(y_test, y_pred3)

#plotting the curves
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve LR (area = {roc_auc})')
plt.plot(fpr2, tpr2, color='green', lw=2, label=f'ROC curve Gridrf (area = {roc_auc2})')
plt.plot(fpr3, tpr3, color='red', lw=2, label=f'ROC curve tree (area = {roc_auc3})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```




##Model Selection

- The approach used in selecting the final model was by comparing the performance between the train and test sets of data.
- First step was looking at the models that showed better generalisation between the train and test set. Meaning the gap between the train and test metrics that exhibited the least difference had the highest priority.
- Metrics that exhibited a range of 90% and above suggests overfitting therefore these models were not considered.
- From the ROC curve above, it shows that The randomforest classifier is performing the best due to its large area under the curve. However this model was not selected because due to the huge difference in the train and test metrics.
- Finally settled on Logistic Regression (Baseline) as the best performing model due to its great generalisation. The metrics in the training and test set also exhibit very minimal difference which is good.

```
In [ ]: # Plotting confusionmatrixDisplay
# Logistic regression
plt.figure(figsize=(8, 6))
cm = confusion_matrix(y_test, y_predr1)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=lr.classes_)
disp.plot(cmap=plt.cm.Blues, values_format='d')
plt.title('Logistic Regression Confusion Matrix')
plt.show();
```

<Figure size 800x600 with 0 Axes>



Observations

- True Negatives (TN): 645 - The model correctly predicted that 645 customers would not use the coupon.
- False Positives (FP): 483 - The model incorrectly predicted that 483 customers would use the coupon when they actually did not.
- False Negatives (FN): 315 - The model incorrectly predicted that 315 customers would not use the coupon when they actually did.
- True Positives (TP): 1094 - The model correctly predicted that 1094 customers would use the coupon.

Analysis of the Results:

- Strengths: The model has a relatively high recall, indicating that it is reasonably good at identifying customers who will use the coupon.
- Limitations: There is still room for improvement. Approximately 22% of customers who used the coupon were missed by the model.

Conclusion

- The Logistic Regression Baseline model provides valuable insights for targeting potential coupon users effectively, helping drive customer engagement. It's a robust initial model, though further refinements could help achieve even greater predictive recall and precision. These insights can support decision-making around targeted marketing strategies, ultimately aiding in maximizing coupon acceptance rates while minimizing unnecessary outreach.